

CS 180 Project 1

Colorizing Prokudin-Gorskii Plates

CS 180 Project 1: Colorizing Prokudin-Gorskii Plates

Brief overview: I take a single tall grayscale scan with three equal-height plates (top→bottom = B, G, R). I split these channels, *align G→B and R→B using integer shifts*, and composite to RGB. For small JPGs I use single-scale exhaustive search; for large TIFs I use a coarse-to-fine pyramid. Only compressed JPG results are shown here, no TIFs.



Detailed Approach

High-level: I split the stacked image into **b**, **g**, and **r** channels, then found 2D integer translations that best align **g** and **r** to **b**. The search window for the single-scale version is $[-R, R] \times [-R, R]$. For TIFs, I built a pyramid by $2\times$ average-pooling. I aligned at the coarsest level and refined at each finer level using appropriately smaller windows.

- **Metrics:**

- **ncc**: normalized cross-correlation on intensity.
- **ssd**: I maximize **-SSD**.
- **edges/edges5**: Sobel magnitude (light Gaussian blur for **edges5**), scored by **-SSD**; robust when exposures differ by channel.

- **auto**: I evaluate `ncc` and `edges5` fairly by scoring both on a comparable scale and pick the better one.
- **Border policy**: during scoring I ignore an inner band (`--border`) to avoid black frames dominating. After shifting, I crop to the common overlap across B/G/R using the *applied* vectors. This removes wrap-around color edges. A small `--pre-crop` helps with thick frames.
- **What I tried when things failed:**
 1. Increase `--max-radius` and confirm via full score grids (debug) if the optimum was clipped.
 2. Switch metric to `edges5` on images with strong exposure differences (e.g., `emir.tif`).
 3. Increase `--border` (or asymmetric left/right) to down-weight thick black frames.
 4. Try a different anchor (`--anchor g`) and compare the score maps.
 5. Use the pyramid path for large shifts in TIFs.



Key Issues & Fixes

1) Shift-sign bug → ghosting

My search returns where the moving channel *is* relative to the reference (`(dy, dx)`). To align it, I **apply** (`-dy, -dx`). Using the wrong sign produced double-ghosting; fixing the sign and reporting *applied* vectors resolved it.

2) `auto` metric fairness (Emir)

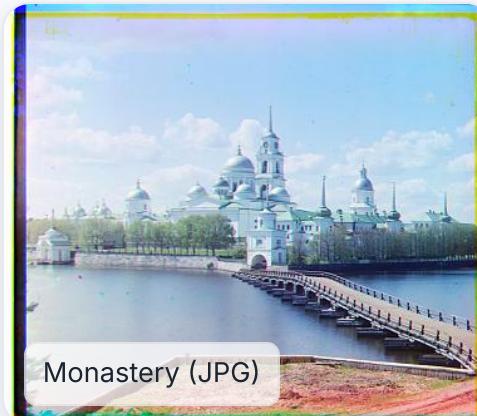
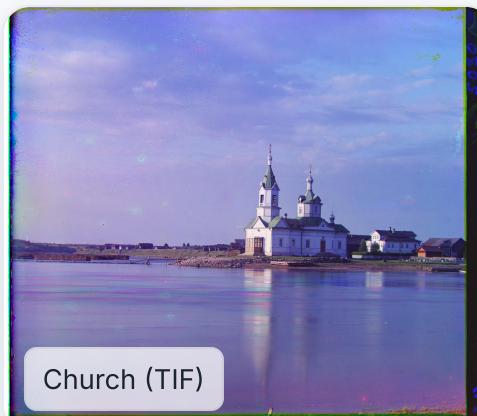
`emir.tif` has different per-channel exposures; intensity NCC can pick a wrong peak. My `auto` compares `ncc` and `edges5` in their own feature spaces but on a comparable numeric scale, so I don't bias toward NCC. For batch runs, I still force `edges5` on TIFs for safety.

3) Frames dominating scores

Thick black borders can dominate SSD/NCC. I ignore an inner band during scoring and crop the final overlap after shifting. This simple policy was more reliable for me than full automatic frame detection.



Results on Provided Images

**single + ncc**Offsets: G→B (5, 2),
R→B (12, 3)**single + ncc**Offsets: G→B (-3, 2),
R→B (3, 2)**single + ncc**Offsets: G→B (3, 3),
R→B (6, 3)**pyramid + edges5**Offsets: G→B
(25, 4), R→B (58,
-4)



Emir (TIF)

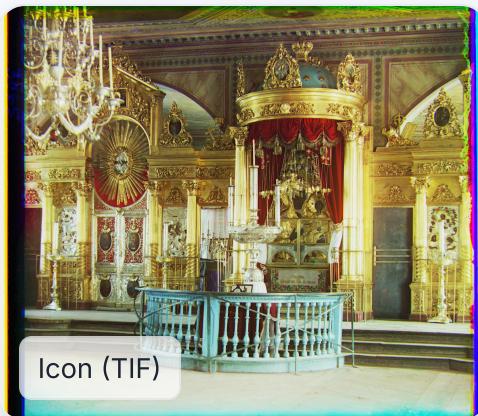


Harvesters (TIF)

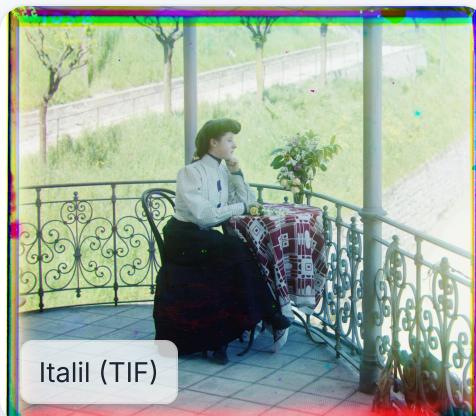
Offsets: G→B

pyramid + edges5 (49, 23), R→B
(107, 40)

Offsets: G→B

pyramid + edges5 (60, 17), R→B
(123, 14)

Icon (TIF)

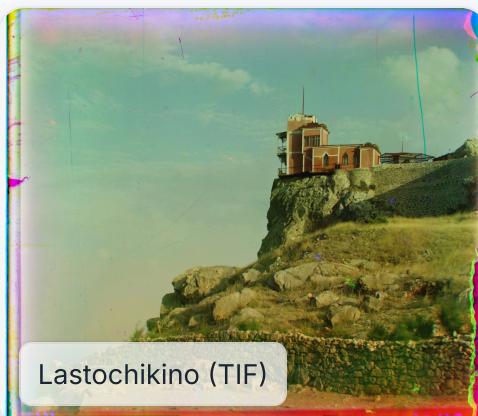


Italil (TIF)

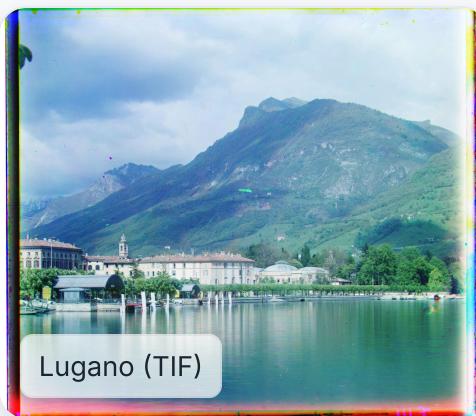
Offsets: G→B

pyramid + edges5 (42, 17), R→B
(90, 23)

Offsets: G→B

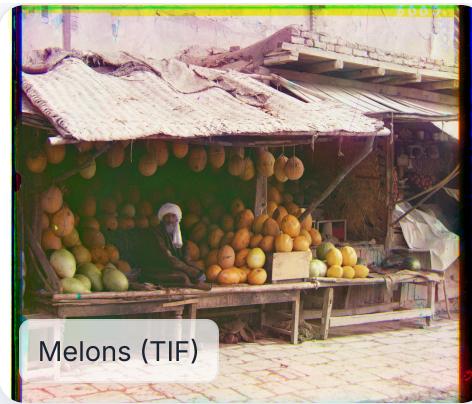
pyramid + edges5 (38, 22), R→B
(77, 36)

Lastochikino (TIF)

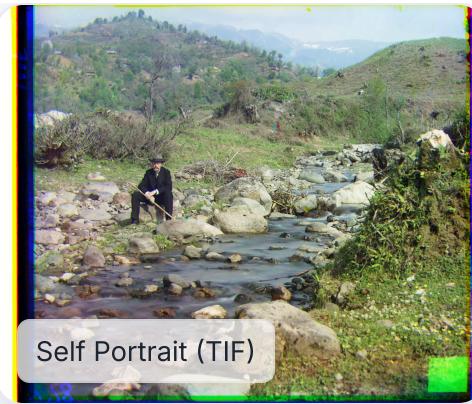


Lugano (TIF)

pyramid + edges5 Offsets: G→B (-3,
-2), R→B (76, -8)**pyramid + edges5** Offsets: G→B (41,
-17), R→B (93,
-29)

**pyramid + edges5**

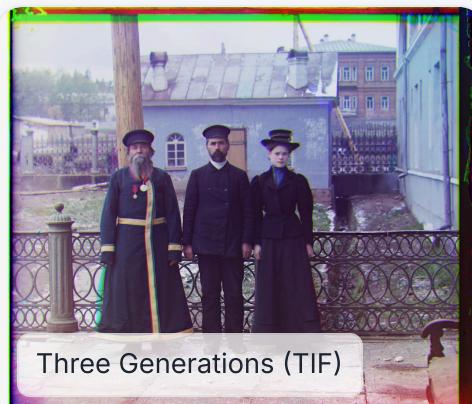
Offsets: G→B (81, 10), R→B (178, 13)

**pyramid + edges5**

Offsets: G→B (79, 30), R→B (176, 37)

**pyramid + edges5**

Offsets: G→B (49, -6), R→B (96, -25)

**pyramid + edges5**

Offsets: G→B (53, 13), R→B (111, 10)



Displacement Vectors (applied)

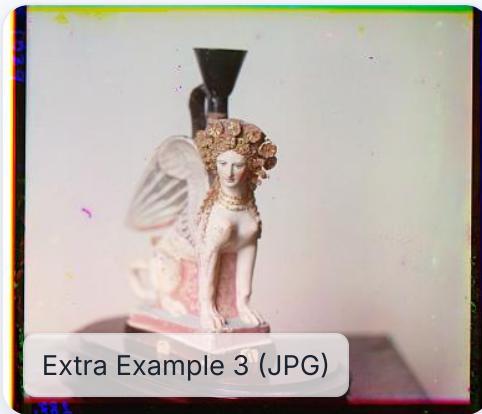
The table lists the *applied* G→B and R→B shifts (dy, dx). To keep the site lightweight, I only include compressed images from [output/](#) (no large TIFs are uploaded).

Image	Method	G→B (dy, dx)	R→B (dy, dx)
Cathedral (JPG)	single + ncc	(5, 2)	(12, 3)
Monastery (JPG)	single + ncc	(-3, 2)	(3, 2)
Tobolsk (JPG)	single + ncc	(3, 3)	(6, 3)
Church (TIF)	pyramid + edges5	(25, 4)	(58, -4)



Image	Method	G→B (dy, dx)	R→B (dy, dx)
Emir (TIF)	pyramid + edges5	(49, 23)	(107, 40)
Harvesters (TIF)	pyramid + edges5	(60, 17)	(123, 14)
Icon (TIF)	pyramid + edges5	(42, 17)	(90, 23)
Italil (TIF)	pyramid + edges5	(38, 22)	(77, 36)
Lastochikino (TIF)	pyramid + edges5	(-3, -2)	(76, -8)
Lugano (TIF)	pyramid + edges5	(41, -17)	(93, -29)
Melons (TIF)	pyramid + edges5	(81, 10)	(178, 13)
Self Portrait (TIF)	pyramid + edges5	(79, 30)	(176, 37)
Siren (TIF)	pyramid + edges5	(49, -6)	(96, -25)
Three Generations (TIF)	pyramid + edges5	(53, 13)	(111, 10)
Extra Example 1 (JPG)	single + ncc	(-3, 1)	(-4, 1)
Extra Example 2 (JPG)	single + ncc	(5, -1)	(10, -3)
Extra Example 3 (JPG)	single + ncc	(2, 3)	(11, 6)

Additional Examples (from the Prokudin-Gorskii collection)

**single + ncc**Offsets: G→B (-3, 1),
R→B (-4, 1)**single + ncc**Offsets: G→B (5, -1),
R→B (10, -3)**single + ncc**Offsets: G→B (2, 3),
R→B (11, 6)

Failures & Difficult Images

- **Emir:** per-channel brightness differs significantly; raw-intensity NCC can mis-score. Using `edges5` (Sobel on a lightly blurred image) focuses on structure and aligns correctly for me.
- **Heavy frames:** when black borders dominate, metrics can chase them. Ignoring an inner band in scoring and cropping the final overlap removed the color bands I often saw at the edges.

How to Reproduce

Single image

- Small JPG (single-scale NCC): `python proj1_colorizer.py --input data/cathedral.jpg --method single --metric ncc --max-radius 15 --auto-crop`
- Large TIF (pyramid + edges): `python proj1_colorizer.py --input data/emir.tif --method pyramid --metric edges5 --auto-crop`



Batch everything

```
# TIFs: pyramid + edges5
for f in data/*.tif; do
    [ -e "$f" ] || continue
    python proj1_colorizer.py --input "$f" --method pyramid
        --output "output/${basename "${f%.*}}_color.jpg"
done

# JPGs: single-scale + NCC
for f in data/*.jpg; do
    [ -e "$f" ] || continue
    python proj1_colorizer.py --input "$f" --method single
        --output "output/${basename "${f%.*}}_color.jpg"
done
```

Notes & Takeaways

- Exhaustive search + clear visualizations helped me debug failures quickly.
- I found edge-based features safer than raw intensity when channels differ in exposure.

- My simplest border policy worked best: ignore during scoring, crop final overlap afterward.
- Future work I'd like to try: subpixel refinement, corner/feature-based matching, and learned border cropping.

© 2025 · CS 180 Project 1

